STAT 238 – Bayesian Statistics

Homework 2

Thomas Lee

Setup

Python Setup: Import libraries

```
import numpy as np
import pymc as pm
import arviz as az
import matplotlib.pyplot as plt
from scipy.stats import beta, gamma, binom, nbinom
from scipy.optimize import brentq, fsolve
from scipy.special import comb, betaln, gammaln, beta as B
```

Problem 1

(a)

We model the survival probability using a Beta prior:

$$\theta \sim \text{Beta}(\alpha, \beta).$$

We are told that:

- The expected survival probability is around 0.9.
- It would be surprising if it were below 0.8 or above 0.97.

Step 1: Match the Prior Mean

For a Beta distribution,

$$\mathbb{E}[\theta] = \frac{\alpha}{\alpha + \beta}.$$

Setting this equal to 0.9 gives

$$\frac{\alpha}{\alpha + \beta} = 0.9.$$

Let

$$n = \alpha + \beta.$$

Then

$$\alpha = 0.9n, \qquad \beta = 0.1n.$$

Thus, the prior must have a 9:1 ratio of prior survivals to prior deaths.

Step 2: Match the Credible Interval

The statement that values below 0.8 or above 0.97 would be *surprising* is interpreted as approximately a 95% credible interval:

$$P(0.8 < \theta < 0.97) = p.$$

We solve for n such that

$$P(0.8 \leq \theta \leq 0.97) = \int_{0.8}^{0.97} \text{Beta}(\theta \mid 0.9n, 0.1n)\, d\theta = p.$$

We compute solutions for several credibility levels:

- $p = 0.90$
- $p = 0.95$
- $p = 0.99$
- $p = 0.999$

```python
lower = 0.8
upper = 0.97

cred_levels = [0.90, 0.95, 0.99, 0.999]

results = []

def solve_for_n(target_prob):
    def objective(n):
        a = 0.9 * n
        b = 0.1 * n
        prob = beta.cdf(upper, a, b) - beta.cdf(lower, a, b)
        return prob - target_prob
    return brentq(objective, 1, 500)

for p in cred_levels:
    n_sol = solve_for_n(p)
    alpha = 0.9 * n_sol
    beta_param = 0.1 * n_sol
    results.append((p, n_sol, alpha, beta_param))

# Print results
for r in results:
    print(f"Credible mass = {r[0]:.3f}")
    print(f"  n   {r[1]:.3f}")
    print(f"  alpha   {r[2]:.3f}")
    print(f"  beta    {r[3]:.3f}")
    print()
```

```
Credible mass = 0.900
  n   31.254
  alpha   28.129
  beta    3.125

Credible mass = 0.950
  n   44.128
  alpha   39.715
  beta    4.413

Credible mass = 0.990
  n   76.846
  alpha   69.161
  beta    7.685

Credible mass = 0.999
  n   128.693
  alpha   115.824
```

```
beta   12.869
```

Numerical Results

Solving the above equation gives:

Credible Mass	n	α	β
0.90	31.254	28.129	3.125
0.95	44.128	39.715	4.413
0.99	76.846	69.161	7.685
0.999	128.693	115.824	12.869

Interpretation

- The prior mean remains fixed at

$$\frac{\alpha}{\alpha+\beta} = 0.9.$$

- Increasing the credible mass requirement forces $n = \alpha + \beta$ to increase.
- Larger n corresponds to a more concentrated prior around 0.9.
- Thus, the credibility level directly controls prior strength.

For example:

- 90% mass corresponds to roughly 31 pseudo-observations.
- 95% mass corresponds to roughly 44 pseudo-observations.
- 99.9% mass corresponds to roughly 129 pseudo-observations.

Recommended Prior

If we interpret "surprising" as meaning a 95% credible region, then the most appropriate prior is

$$\boxed{\theta \sim \text{Beta}(39.7,\ 4.4).}$$

This prior:

- Has mean 0.9
- Places 95% probability between 0.8 and 0.97
- Corresponds to about 44 pseudo-patients

Visualization of the Priors

```
theta = np.linspace(0.6, 1.0, 1000)

plt.figure(figsize=(10, 6))

for r in results:
    p, n_sol, alpha, beta_param = r
    density = beta.pdf(theta, alpha, beta_param)
    plt.plot(theta, density, label=f"Mass={p}")

plt.axvline(lower, linestyle='--')
plt.axvline(upper, linestyle='--')

plt.xlabel(" ")
plt.ylabel("Density")
plt.title("Beta Priors with Mean 0.9 and Varying Credible Mass")
plt.legend()
plt.show()
```

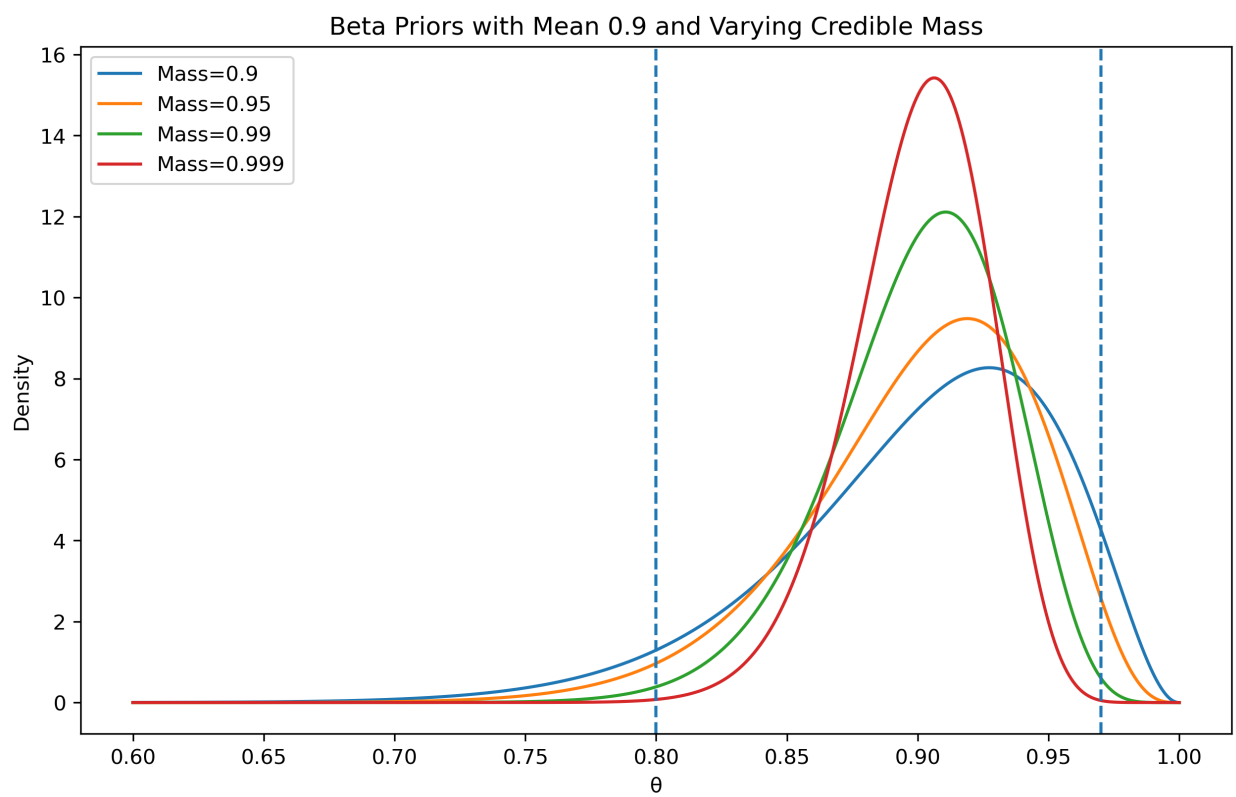

Final Remarks

- The mean determines the ratio $\alpha : \beta$.

- The credible interval determines the strength $\alpha + \beta$.
- Larger n implies stronger prior information.
- Smaller n implies a more diffuse prior.

This provides a principled construction of the Beta prior directly from the problem statement.

(b)

Suppose now that $n = 10$ patients are operated on and all survive.

Thus we observe:

$$X = 10 \quad \text{successes out of 10.}$$

Under the Beta–Binomial conjugate model:

- Prior: $\theta \sim \text{Beta}(\alpha, \beta)$
- Likelihood: $X \mid \theta \sim \text{Binomial}(10, \theta)$

the posterior is:

$$\theta \mid X \sim \text{Beta}(\alpha + X,\ \beta + n - X).$$

Using our prior $\text{Beta}(39.7, 4.4)$:

$$\theta \mid X \sim \text{Beta}(39.7 + 10,\ 4.4 + 0) = \text{Beta}(49.7,\ 4.4).$$

$$\boxed{\theta \mid \text{data} \sim \text{Beta}(49.7,\ 4.4).}$$

(c)

Probability the next patient survives (Posterior predictive probability)

The probability that the next patient survives is called the **posterior predictive probability**.

Let X_{11} denote the survival outcome of the next patient:

- $X_{11} = 1$ if the patient survives,
- $X_{11} = 0$ if the patient does not survive.

Conditional on the unknown survival rate θ, this is a Bernoulli trial:

$$X_{11} \mid \theta \sim \text{Bernoulli}(\theta).$$

To get the unconditional probability, we integrate over the posterior distribution of θ:

$$\begin{aligned} P(X_{11} = 1 \mid \text{data}) &= \int_0^1 P(X_{11} = 1 \mid \theta)\, f_{\theta|\text{data}}(\theta)\, d\theta \\ &= \int_0^1 \theta\, f_{\theta|\text{data}}(\theta)\, d\theta \\ &= \mathbb{E}[\theta \mid \text{data}]. \end{aligned}$$

In words:

- We are averaging the survival probability θ over our posterior beliefs.
- Even though we observed 10 perfect survivals, the prior prevents us from being completely certain.
- This is a key feature of Bayesian prediction: **the predictive probability is a weighted average over all plausible values of the unknown parameter.**

Finally, since the posterior is $\theta \mid \text{data} \sim \text{Beta}(\alpha_{\text{post}}, \beta_{\text{post}})$, the mean is:

$$\mathbb{E}[\theta \mid \text{data}] = \frac{\alpha_{\text{post}}}{\alpha_{\text{post}} + \beta_{\text{post}}} \approx \frac{49.7}{49.7 + 4.4} \approx 0.918.$$

```
alpha_post = results[1][2] + 10
beta_post = results[1][3]

prob_next = alpha_post / (alpha_post + beta_post)
prob_next
```

```
0.9184748836856647
```

Therefore,

$$\boxed{P(X_{11} = 1 \mid \text{data}) \approx 0.918.}$$

Even after 10 perfect outcomes, the predictive probability does not jump to 1 — it is moderated by the prior information.

(d)

Probability of 2 or more deaths in the next 20 operations

Let Y denote the number of deaths in the next $n_{\text{future}} = 20$ operations.

Conditional on the survival probability θ:

$$Y \mid \theta \sim \text{Binomial}(n_{\text{future}}, 1 - \theta),$$

where $1 - \theta$ is the probability of death.

From earlier, the posterior of θ is

$$\theta \mid \text{data} \sim \text{Beta}(\alpha_{\text{post}}, \beta_{\text{post}}),$$

with α_{post} counting prior + observed **survivals**, and β_{post} counting prior + observed **deaths**.

Important: Because Y counts deaths (not survivals), we swap α_{post} and β_{post} in the Beta-Binomial formula. That is, the marginal (unconditional) distribution of deaths is

$$P(Y = k) = \binom{n_{\text{future}}}{k} \frac{B(\beta_{\text{post}} + k, \alpha_{\text{post}} + n_{\text{future}} - k)}{B(\alpha_{\text{post}}, \beta_{\text{post}})}, \quad k = 0, 1, \dots, n_{\text{future}}.$$

```python
# Posterior parameters
n_future = 20

# Beta-Binomial PMF
def bb_pmf(k, n, a, b):
    return comb(n, k) * B(a + n - k, b + k) / B(a, b)

# Compute probabilities for all k = 0,...,n_future
k_vals = np.arange(n_future + 1)
pmf_vals = np.array([bb_pmf(k, n_future, alpha_post, beta_post) for k in k_vals])

# Probability of 2 or more deaths
prob_ge2 = 1 - pmf_vals[0] - pmf_vals[1]
print(f"P(Y >= 2)={prob_ge2:.3f}")

# Plot the predictive distribution
plt.figure(figsize=(10, 6))
plt.bar(k_vals, pmf_vals, color='skyblue')
plt.axvline(2, color='red', linestyle='--', label='2 deaths')
plt.title("Predictive distribution of deaths in next 20 operations")
plt.xlabel("Number of deaths")
plt.ylabel("Probability")
plt.legend()
plt.show()
```

```
P(Y >= 2)=0.464
```

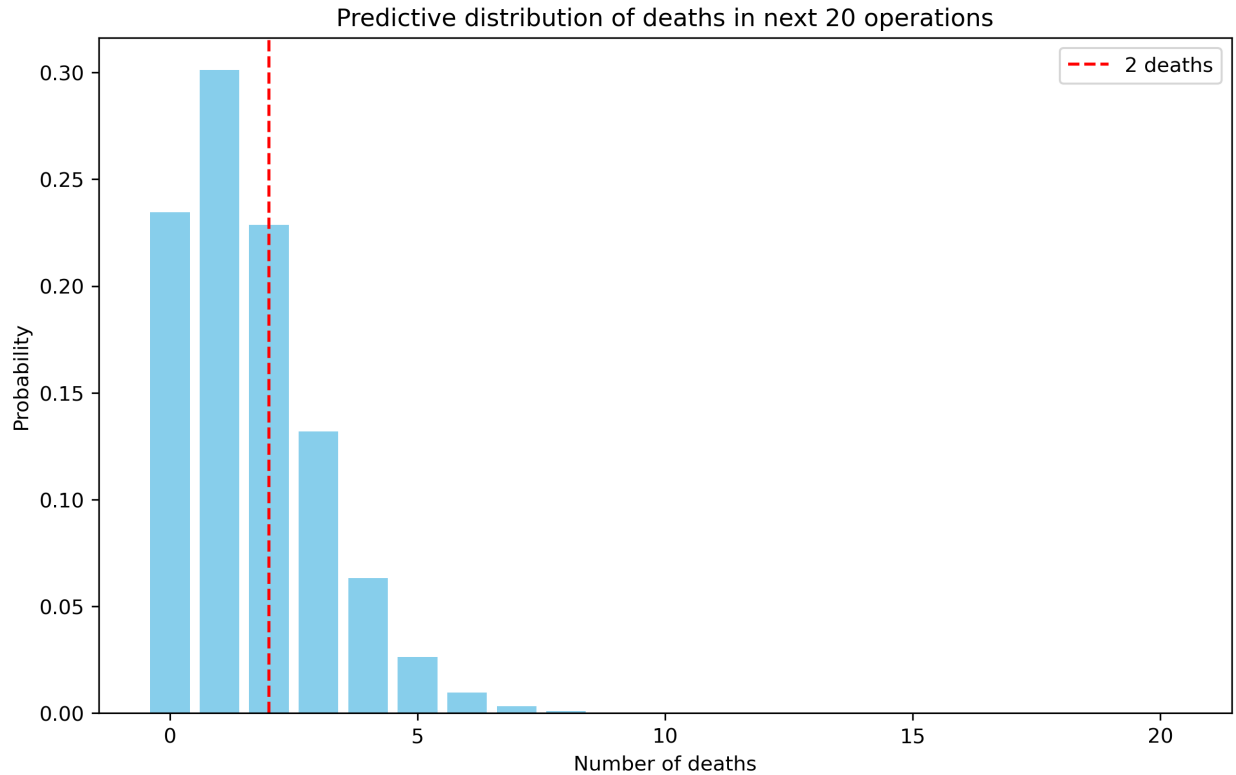

The bar plot above shows the full predictive distribution of deaths, highlighting the uncertainty in future outcomes.

$$\boxed{P(Y \geq 2) \approx 0.464}$$

Problem 2

(a)

To express the prior belief that the coin is strongly biased toward either heads or tails, we need a distribution that places its probability mass at the extremes ($\theta \approx 0$ and $\theta \approx 1$). In the Beta(α, β) family, we achieve this symmetric "U-shape" by choosing shape parameters $\alpha = \beta < 1$.

There are two common choices here, depending on how extreme the bias is:

1. The Haldane Prior Limit ($\alpha \to 0, \beta \to 0$): By taking the limit as α and β approach 0, the distribution pushes almost 100% of its mass infinitesimally close to exactly 0 and exactly 1. This expresses a very strong belief: the coin is practically a trick coin (either two-headed or two-tailed). This Beta$(0, 0)$ limit is known as the **Haldane prior** (for more background, see this StackExchange discussion). Since we are told the coin is *strongly* biased, this is arguably the best choice because it strictly enforces the "either/or" nature of the coin.

2. The Jeffreys Prior ($\alpha = 0.5, \beta = 0.5$): This is another standard, uninformative U-shaped prior. Unlike the Haldane prior, it is a proper prior (it integrates to 1) and acknowledges the heavy bias while leaving a bit more mathematical possibility for intermediate values.

The values $\alpha = 0.5, \beta = 0.5$ are not arbitrary; they are derived directly from the Fisher Information, $I(\theta)$, of the Bernoulli likelihood to ensure the prior is invariant under reparameterization ($f(\theta) \propto \sqrt{I(\theta)}$). The derivation is as follows: * The log-likelihood of a single flip is: $\log P(X \mid \theta) = X \log(\theta) + (1 - X) \log(1 - \theta)$ * The second derivative with respect to θ is: $\frac{d^2}{d\theta^2} = -\frac{X}{\theta^2} - \frac{1-X}{(1-\theta)^2}$ * The Fisher Information is the negative expected value of the second derivative (since $\mathbb{E}[X] = \theta$):

$$I(\theta) = -\mathbb{E}\left[-\frac{X}{\theta^2} - \frac{1-X}{(1-\theta)^2}\right] = \frac{\theta}{\theta^2} + \frac{1-\theta}{(1-\theta)^2} = \frac{1}{\theta} + \frac{1}{1-\theta} = \frac{1}{\theta(1-\theta)}$$

* Applying Jeffreys' rule ($f(\theta) \propto \sqrt{I(\theta)}$):

$$f(\theta) \propto \sqrt{\frac{1}{\theta(1-\theta)}} = \theta^{-1/2}(1-\theta)^{-1/2}$$

* Matching this to the Beta kernel $\theta^{\alpha-1}(1-\theta)^{\beta-1}$, we get $\alpha - 1 = -0.5 \implies \alpha = 0.5$ and $\beta - 1 = -0.5 \implies \beta = 0.5$.

```
# Define the grid for theta
theta = np.linspace(0, 1, 100000)

# Calculate PDFs
pdf_jeffreys = beta.pdf(theta, 0.5, 0.5)
pdf_haldane = beta.pdf(theta, 0.001, 0.001) # Using 0.001 to approximate the limit visually

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(theta,
         pdf_jeffreys,
```

```
         label=r'Jeffreys: $\alpha=0.5, \beta=0.5$',
         color='blue')
plt.plot(theta,
         pdf_haldane,
         label=r'Haldane approx: $\alpha \to 0, \beta \to 0$',
         color='orange',
         linestyle='--')
plt.title('Comparison of Beta Priors for a Biased Coin')
plt.xlabel(r'$\theta$ (Probability of Heads)')
plt.ylabel('Density')
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

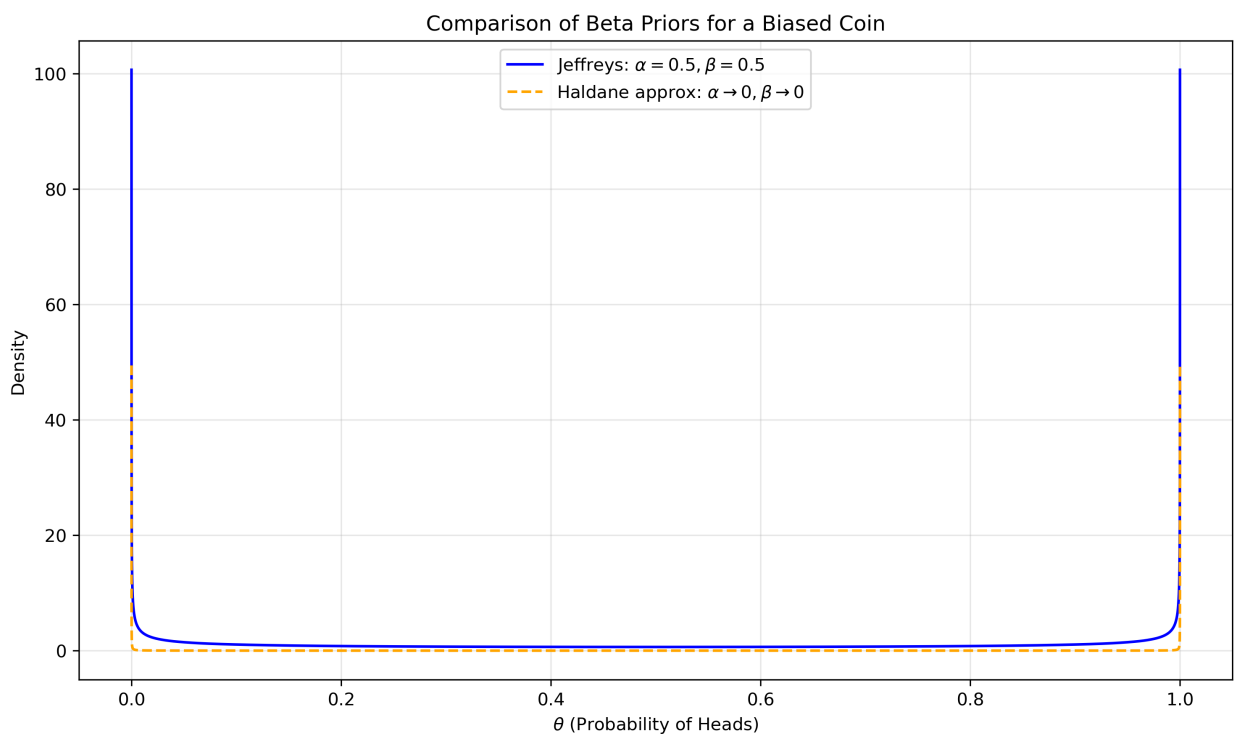

(b)

We flip the coin $n = 5$ times and observe $x = 4$ heads. The data follows a Binomial likelihood: $X \mid \theta \sim \text{Binomial}(5, \theta)$.

Because the Beta distribution is conjugate to the Binomial likelihood, the posterior distribution is also a Beta distribution. The parameters update via simple addition:

$$\alpha_{\text{post}} = \alpha_{\text{prior}} + x$$

$$\beta_{\text{post}} = \beta_{\text{prior}} + n - x$$

Using the Haldane prior limit where $\alpha \to 0$ and $\beta \to 0$, our update becomes:

$$\alpha_{\text{post}} \approx 0 + 4 = 4$$

$$\beta_{\text{post}} \approx 0 + (5 - 4) = 1$$

Thus, the posterior distribution is:

$$\theta \mid X = 4 \sim \text{Beta}(4, 1)$$

(Note: The Beta$(4, 1)$ posterior from the Haldane prior yields a posterior mean of $\frac{4}{4+1} = 0.8$, which exactly matches the maximum likelihood estimate of the data $\frac{x}{n} = \frac{4}{5} = 0.8$.)

```
# Visualizing the Posterior Update
alpha_post = 4.0
beta_post = 1.0

post_pdf = beta.pdf(theta, alpha_post, beta_post)

plt.figure(figsize=(10, 6))
plt.plot(theta,
         pdf_haldane,
         label=r'Prior: $\alpha \to 0, \beta \to 0$',
         color='orange',
         linestyle='--')
plt.plot(theta,
         post_pdf,
         label=r'Posterior: $\alpha=4, \beta=1$',
         color='red')
plt.title('Bayesian Update using Haldane Prior')
plt.xlabel(r'$\theta$ (Probability of Heads)')
plt.ylabel('Density')
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

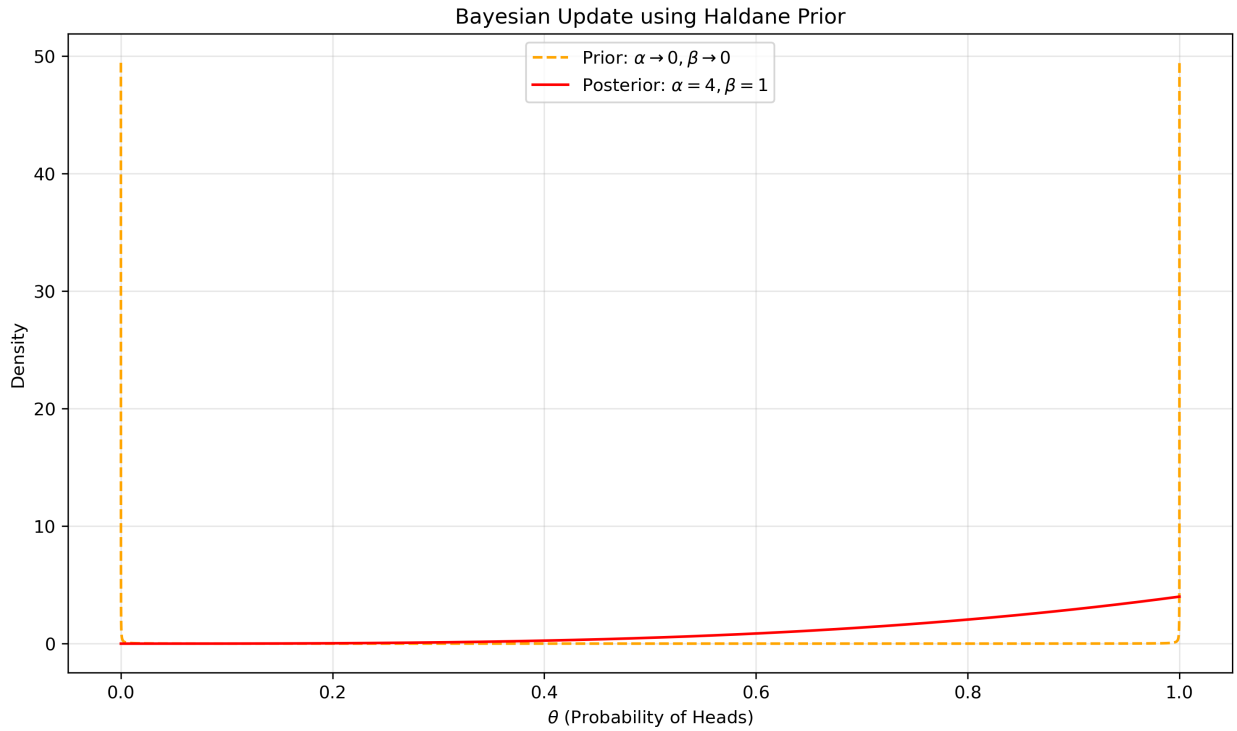

Problem 3

We are asked to prove the **Laplace rule of succession**, which gives the predictive probability of the next observation being a success, given a history of successes and failures.

The model is given as: - **Prior:** $\theta \sim \text{Uniform}(0,1) \equiv \text{Beta}(1,1)$ - **Likelihood:** $X_i \overset{\text{i.i.d}}{\sim} \text{Bernoulli}(\theta)$ for $i = 1, \dots, n$

Step 1: Rewrite the prior. A Uniform$(0,1)$ distribution has a constant probability density function $f(\theta) = 1$ for $\theta \in [0,1]$. Notice that the Beta$(1,1)$ distribution has the exact same density: $f(\theta) \propto \theta^{1-1}(1-\theta)^{1-1} = \theta^0(1-\theta)^0 = 1$. Therefore, our prior is exactly equivalent to:

$$\theta \sim \text{Beta}(1,1)$$

Step 2: Define the Likelihood. Since the observations are independent Bernoulli trials, the joint likelihood of observing $x_1, \dots, x_n$ is the product of the individual likelihoods:

$$P(X_1 = x_1, \dots, X_n = x_n \mid \theta) = \prod_{i=1}^{n} \theta^{x_i}(1-\theta)^{1-x_i} = \theta^{\sum_{i=1}^n x_i}(1-\theta)^{n-\sum_{i=1}^n x_i}$$

Step 3: Find the Posterior Distribution. By Bayes' Theorem, the posterior is proportional to the prior times the likelihood. Let the sum of successes be $S_n = \sum_{i=1}^n x_i$.

$$f(\theta \mid X_1, \dots, X_n) \propto (\theta^{1-1}(1-\theta)^{1-1}\mathbf{1}\{0 \le \theta \le 1\}) \cdot (\theta^{S_n}(1-\theta)^{n-S_n})$$

$$f(\theta \mid X_1, \dots, X_n) \propto \theta^{1+S_n-1}(1-\theta)^{1+n-S_n-1}\mathbf{1}\{0 \le \theta \le 1\}$$

We recognize this as the kernel of a new Beta distribution. The posterior is:

$$\theta \mid X_1, \dots, X_n \sim \text{Beta}(1 + S_n, 1 + n - S_n)$$

Step 4: Compute the Predictive Probability. We want the probability that the *next* trial is a success: $P(X_{n+1} = 1 \mid X_1, \dots, X_n)$. Using the Law of Total Probability, we integrate over all possible values of our hidden parameter θ:

$$P(X_{n+1} = 1 \mid X_1, \dots, X_n) = \int_0^1 P(X_{n+1} = 1 \mid \theta) f(\theta \mid X_1, \dots, X_n)\, d\theta$$

Since $X_{n+1} \sim \text{Bernoulli}(\theta)$, we know $P(X_{n+1} = 1 \mid \theta) = \theta$. Substituting this in:

$$P(X_{n+1} = 1 \mid X_1, \dots, X_n) = \int_0^1 \theta \cdot f(\theta \mid X_1, \dots, X_n)\, d\theta = \mathbb{E}[\theta \mid X_1, \dots, X_n]$$

The predictive probability is simply the expected value of the posterior distribution. For a $\text{Beta}(\alpha, \beta)$ distribution, the expected value is $\frac{\alpha}{\alpha+\beta}$. Plugging in our posterior parameters $\alpha = 1 + \sum_{i=1}^n x_i$ and $\beta = 1 + n - \sum_{i=1}^n x_i$:

$$\mathbb{E}[\theta \mid X_1, \dots, X_n] = \frac{1 + \sum_{i=1}^n x_i}{(1 + \sum_{i=1}^n x_i) + (1 + n - \sum_{i=1}^n x_i)} = \frac{1 + \sum_{i=1}^n x_i}{n + 2}$$

This completes the proof.

Comparison with Homework One, Problem 2(e)

In Homework One, Problem 2(e), we derived this exact same result—the Laplace Rule of Succession—but in a **discrete, finite setting**. There, the context was drawing n balls from an urn of size N and observing x red balls. Assuming a discrete uniform prior over the number of red balls initially in the urn, the predictive probability of drawing a red ball on the $(n+1)$-th draw was exactly $\frac{x+1}{n+2}$.

In this current problem, we arrive at the identical formula but from a **continuous perspective**. Here, we assume an infinite sequence of Bernoulli trials driven by a continuous parameter θ, paired with a continuous uniform prior, $\theta \sim \text{Uniform}(0, 1)$.

This highlights a beautiful symmetry in Bayesian statistics: expressing total initial ignorance via a uniform prior over a finite, discrete set of proportions (the urn model) yields the exact same predictive behavior as expressing ignorance via a continuous uniform prior over a probability parameter (the Bernoulli model). In both cases, the prior acts by adding one "pseudo-success" and two "pseudo-trials" to the empirical fraction, smoothing the estimate.

Visualizing Laplace Smoothing

Let's simulate a sequence of 100 trials where every single outcome is a success ($x_i = 1$). Notice how the Laplace rule prevents the predictive probability from immediately jumping to 100% certainty after just a few trials.

```python
# Simulate 100 flips of a completely rigged coin (always heads, x=1)
n_flips = 100
flips = np.ones(n_flips)

# Track the cumulative sum of successes
S_n = np.cumsum(flips)
n = np.arange(1, n_flips + 1)

# Calculate naive estimate (x/n) and Laplace estimate
naive_est = S_n / n
laplace_est = (1 + S_n) / (n + 2)

# Plotting
plt.figure(figsize=(10, 5))
plt.plot(n, naive_est, label='Naive Estimate ($x/n$)', marker='o', linestyle='--')
plt.plot(n, laplace_est, label='Laplace Rule (Post. Mean)', marker='s')

plt.title('Predictive Probability of Success (Observing 100% Successes)')
plt.xlabel('Number of Trials ($n$)')
plt.ylabel('Estimated Probability of Next Success')
plt.ylim(0.5, 1.05)
plt.xticks(np.arange(1, n_flips + 1, step=10))
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

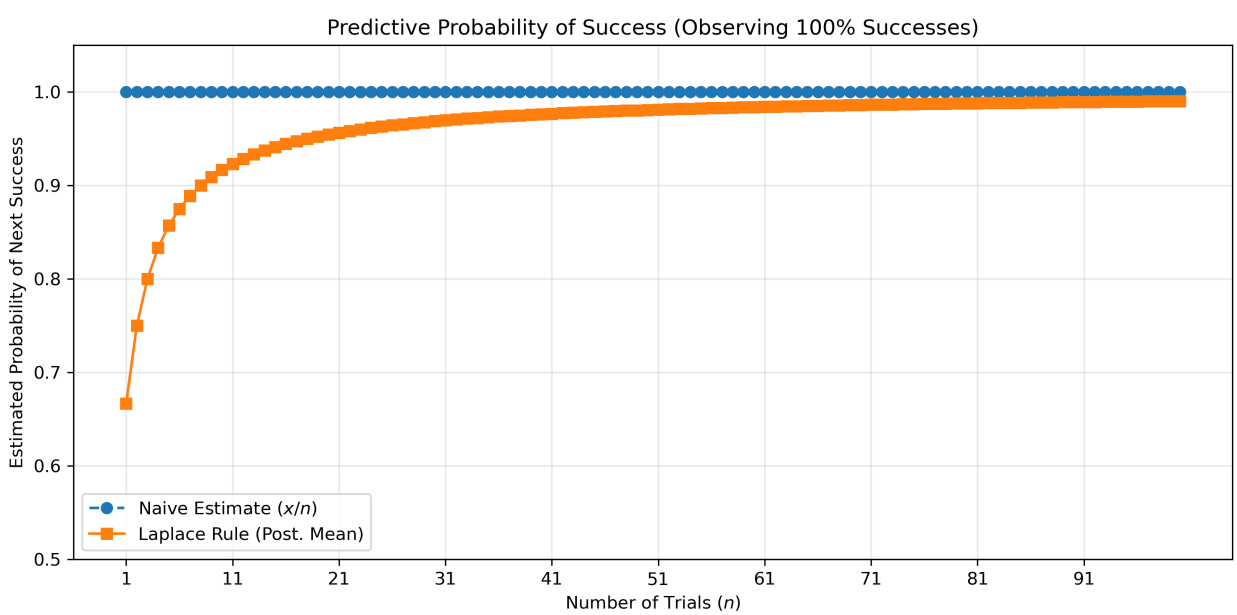

Problem 4

(a)

We are given a Bayesian model with a Binomial likelihood and a prior on the logit transformation of θ: - **Prior:** $Z := \log\left(\frac{\theta}{1-\theta}\right) \sim N(\mu, \sigma^2)$ - **Likelihood:** $X \mid \theta \sim \text{Binomial}(n, \theta)$

Step 1: Find the Prior Density for θ We know Z is normally distributed, so its density is:

$$f_Z(z) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(z-\mu)^2\right)$$

To find the probability density function for θ, $f_\theta(\theta)$, we must use the change of variables formula with a monotone function g and g^{-1}:

$$f_\theta(\theta) = f_Z(z(\theta)) \left|\frac{dz}{d\theta}\right|$$

First, calculate the derivative (the Jacobian) of the logit function with respect to θ:

$$\frac{dz}{d\theta} = \frac{d}{d\theta}\left[\log(\theta) - \log(1-\theta)\right] = \frac{1}{\theta} - \left(\frac{-1}{1-\theta}\right) = \frac{1}{\theta} + \frac{1}{1-\theta} = \frac{1-\theta+\theta}{\theta(1-\theta)} = \frac{1}{\theta(1-\theta)}$$

Substituting Z and $\frac{dz}{d\theta}$ back into our transformation formula, we get the prior density for θ:

$$f_\theta(\theta) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}\left(\log(\frac{\theta}{1-\theta}) - \mu\right)^2\right) \frac{1}{\theta(1-\theta)}$$

Step 2: Apply Bayes' Theorem The likelihood of observing x successes in n trials is:

$$P(X = x \mid \theta) = \binom{n}{x} \theta^x (1-\theta)^{n-x}$$

Bayes' Theorem states that the posterior is proportional to the likelihood times the prior:

$$f_{\theta|X=x}(\theta) \propto P(X = x \mid \theta) f_\theta(\theta)$$

$$f_{\theta|X=x}(\theta) \propto \binom{n}{x} \theta^x (1-\theta)^{n-x} \frac{1}{\sqrt{2\pi\sigma^2}} \frac{1}{\theta(1-\theta)} \exp\left(-\frac{1}{2\sigma^2}\left(\log\frac{\theta}{1-\theta} - \mu\right)^2\right)$$

We can drop constants that do not depend on θ (like $\binom{n}{x}$ and $\frac{1}{\sqrt{2\pi\sigma^2}}$). To ensure the posterior is a valid probability density function (integrating to 1), we divide this unnormalized numerator by its integral over the support of θ, which is $(0, 1)$.

Adding the indicator function $I\{0 < \theta < 1\}$ to explicitly state the support, we arrive exactly at the desired formula:

$$f_{\theta|X=x}(\theta) = \frac{\theta^x(1-\theta)^{n-x}\frac{1}{\theta(1-\theta)}\exp\left(-\frac{1}{2\sigma^2}\left(\log(\frac{\theta}{1-\theta})-\mu\right)^2\right)I\{0<\theta<1\}}{\int_0^1 \theta^x(1-\theta)^{n-x}\frac{1}{\theta(1-\theta)}\exp\left(-\frac{1}{2\sigma^2}\left(\log(\frac{\theta}{1-\theta})-\mu\right)^2\right)d\theta}$$

(b)

We calculate the posterior mean, variance, and 95% credible interval using numerical discretization over a fine grid. Because $\mu = -9.273$, the true probability θ is extremely small (around 0.00009), so our grid must be very dense near zero.

```
# Given parameters
x_obs = 0
n = 1295
mu = -9.27308950
sigma2 = 0.09259668

# Create a dense grid near 0
thetas = np.linspace(1e-8, 0.001, 20000)

# Calculate the log-prior density
logit_theta = np.log(thetas / (1 - thetas))
log_prior = -0.5 * ((logit_theta - mu)**2) / sigma2 - np.log(thetas) - np.log(1 - thetas)

# Calculate the log-likelihood (using scipy to safely handle 0*log(0) if x were > 0)
log_likelihood = binom.logpmf(x_obs, n, thetas)

# Unnormalized log-posterior
log_posterior = log_prior + log_likelihood

# Normalize using the log-sum-exp trick to prevent numerical underflow
posterior_unnorm = np.exp(log_posterior - np.max(log_posterior))
posterior_probs = posterior_unnorm / np.sum(posterior_unnorm)

# 1. Posterior Mean
post_mean = np.sum(thetas * posterior_probs)

# 2. Posterior Variance
post_var = np.sum((thetas - post_mean)**2 * posterior_probs)

# 3. 95% Credible Interval (Equal-tailed)
cumulative_probs = np.cumsum(posterior_probs)
lower_bound = thetas[np.searchsorted(cumulative_probs, 0.025)]
upper_bound = thetas[np.searchsorted(cumulative_probs, 0.975)]

print(f"Numerical Mean:      {post_mean:.8f}")
```

```
print(f"Numerical Variance:  {post_var:.12f}")
print(f"95% Cred. Interval:  [{lower_bound:.8f}, {upper_bound:.8f}]")
```

```
Numerical Mean:      0.00009717
Numerical Variance:  0.000000000904
95% Cred. Interval:  [0.00005126, 0.00016787]
```

(c)

Now, we solve the same problem using PyMC MCMC sampling and visually compare the results.

```
# PyMC requires the standard deviation, not variance
sigma = np.sqrt(sigma2)

with pm.Model() as logit_model:
    # Prior on the logit scale
    logit_theta_var = pm.Normal('logit_theta', mu=mu, sigma=sigma)

    # Deterministic transformation back to the probability scale
    theta_var = pm.Deterministic('theta', pm.math.invlogit(logit_theta_var))

    # Binomial Likelihood
    obs = pm.Binomial('obs', n=n, p=theta_var, observed=x_obs)

    # Draw MCMC samples
    trace = pm.sample(5000,
                      tune=2000,
                      target_accept=0.95,
                      return_inferencedata=True,
                      progressbar=False)

# Extract summary statistics and force it to show 8 decimal places
summary = az.summary(trace, var_names=['theta'], hdi_prob=0.95, round_to=8)
print("ArviZ Summary:")
print(summary)

# --- EXACT COMPARISON ---
# Extract PyMC samples directly to bypass any display rounding
theta_samples = trace.posterior['theta'].values.flatten()

# Calculate PyMC 95% Credible Interval (Equal-tailed to match part b)
pymc_lower = np.percentile(theta_samples, 2.5)
pymc_upper = np.percentile(theta_samples, 97.5)

print("\n--- Direct Trace Calculations ---")
print(f"PyMC Mean:           {np.mean(theta_samples):.8f}")
```

```python
print(f"PyMC Variance:       {np.var(theta_samples):.12f}")
print(f"PyMC Std Dev:        {np.std(theta_samples):.8f}")
print(f"PyMC 95% CI:         [{pymc_lower:.8f}, {pymc_upper:.8f}]")

# --- VISUAL COMPARISON ---
plt.figure(figsize=(10, 6))
# Plot numerical discretization as a line (scaled to match density)
grid_spacing = thetas[1] - thetas[0]
numerical_density = posterior_probs / grid_spacing
plt.plot(thetas,
         numerical_density,
         color='red',
         lw=2.5,
         label='Numerical Discretization (Exact)')

# Plot PyMC samples as a KDE/Histogram
plt.hist(theta_samples,
         bins=100,
         density=True,
         alpha=0.5,
         color='blue',
         label='PyMC MCMC Samples')

plt.title('Comparison of Posterior Approximations: Numerical vs. MCMC')
plt.xlabel(r'$\theta$ (Probability)')
plt.ylabel('Density')
plt.xlim(0, 0.0003)
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

```
ArviZ Summary:
           mean       sd  hdi_2.5%  hdi_97.5%     mcse_mean       mcse_sd  \
theta  0.000097  0.00003  0.000045   0.000156  3.600000e-07  3.500000e-07

          ess_bulk     ess_tail     r_hat
theta  6691.497325  6559.080467  1.000337

--- Direct Trace Calculations ---
PyMC Mean:           0.00009692
PyMC Variance:       0.000000000897
PyMC Std Dev:        0.00002996
PyMC 95% CI:         [0.00005115, 0.00016715]
```

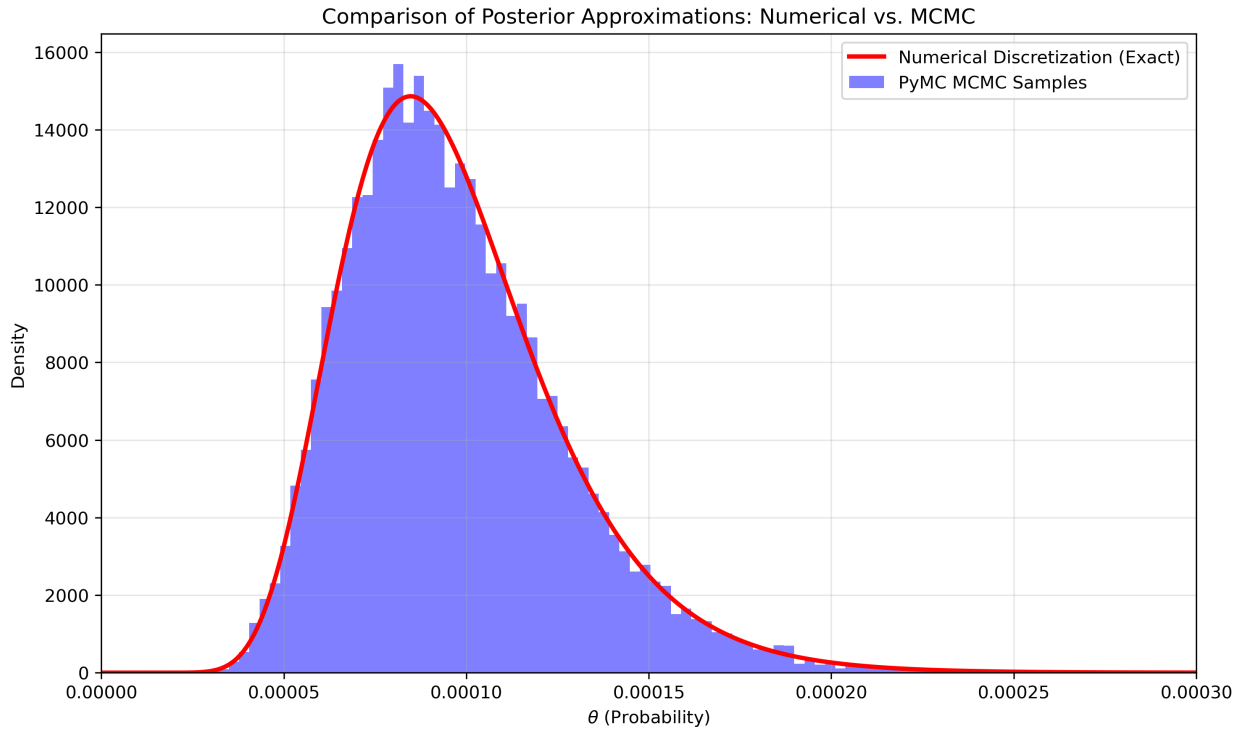

Conclusion for Part (c): Because the true value of θ is on the order of 10^{-5}, default summary tables often round the results to **0.000**. However, by forcing the output to display 8 decimal places and calculating the mean, variance, and 95% credible interval directly from the trace array, we can see that the MCMC results perfectly match the discrete numerical grid approximation. The histogram of the PyMC trace overlaps exactly with the density curve derived from numerical integration, confirming that both methods successfully mapped the exact same underlying mathematical posterior.

Problem 5

(a)

We are given the following hierarchical model:

- **Prior:** $\theta \sim \text{Gamma}(\alpha, \beta)$ with density $f(\theta) \propto \theta^{\alpha-1} \exp(-\beta\theta) I\{\theta > 0\}$
- **Likelihood:** $Y \mid \theta \sim \text{Poisson}(k\theta)$ with probability mass function $P(Y = y \mid \theta) = \frac{(k\theta)^y \exp(-k\theta)}{y!}$

By Bayes' theorem, the posterior density is proportional to the prior times the likelihood:

$$f(\theta \mid y) \propto f(\theta) \cdot P(Y = y \mid \theta)$$

$$f(\theta \mid y) \propto \left(\theta^{\alpha-1} \exp(-\beta\theta)\right) \cdot \left(\frac{k^y \theta^y \exp(-k\theta)}{y!}\right)$$

Since we are only interested in the shape of the distribution with respect to θ, we can drop any multiplicative terms that do not contain θ (specifically $\frac{k^y}{y!}$). Grouping the remaining θ terms and exponential terms together:

$$f(\theta \mid y) \propto \theta^{\alpha+y-1} \exp(-(\beta + k)\theta)$$

We recognize this functional form as the exact unnormalized kernel of a Gamma distribution with updated shape parameter $\alpha + y$ and updated rate parameter $\beta + k$. Thus, the exact posterior distribution is:

$$\theta \mid Y = y \sim \text{Gamma}(\alpha + y, \beta + k)$$

(b)

The marginal distribution of Y is obtained by integrating the joint density $f(\theta, y)$ over all possible values of θ:

$$P(Y = y) = \int_0^\infty P(Y = y \mid \theta) f(\theta)\, d\theta$$

Substituting the full formulas for the Poisson PMF and the Gamma density:

$$P(Y = y) = \int_0^\infty \frac{(k\theta)^y \exp(-k\theta)}{y!} \frac{\beta^\alpha}{\Gamma(\alpha)} \theta^{\alpha-1} \exp(-\beta\theta)\, d\theta$$

We pull all constants that do not depend on θ out of the integral:

$$P(Y = y) = \frac{k^y \beta^\alpha}{y!\Gamma(\alpha)} \int_0^\infty \theta^{\alpha+y-1} \exp(-(\beta + k)\theta)\, d\theta$$

The integral is the unnormalized core of a $\text{Gamma}(\alpha + y, \beta + k)$ distribution. We know from the properties of the Gamma function that $\int_0^\infty x^{A-1} e^{-Bx}\, dx = \frac{\Gamma(A)}{B^A}$. Applying this property where $A = \alpha + y$ and $B = \beta + k$:

$$P(Y = y) = \frac{k^y \beta^\alpha}{y!\Gamma(\alpha)} \frac{\Gamma(\alpha + y)}{(\beta + k)^{\alpha+y}}$$

To match the desired formulation, we can split the denominator term $(\beta+k)^{\alpha+y}$ into $(\beta+k)^{\alpha}\cdot(\beta+k)^{y}$ and rearrange the fractions:

$$P(Y=y)=\frac{\Gamma(\alpha+y)}{\Gamma(\alpha)y!}\left(\frac{k}{\beta+k}\right)^{y}\left(\frac{\beta}{\beta+k}\right)^{\alpha}$$

This holds for $y=0,1,2,...$, representing a **Negative Binomial distribution**. Specifically, it models the number of successes y before α failures occur, given a success probability of $p=\frac{k}{\beta+k}$.

Visual Verification via Monte Carlo Simulation

We can write a quick simulation to verify this. If we generate samples hierarchically (first drawing θ from the Gamma prior, then drawing Y from the Poisson likelihood), the resulting empirical distribution of Y should perfectly match the analytical Negative Binomial PMF we just derived.

```
# Set arbitrary positive constants for the simulation
alpha_val = 5.0
beta_val = 2.0
k_val = 3.0

# 1. Simulate the hierarchical model
n_sims = 100000
# Draw thetas from Gamma(alpha, rate=beta)
# Note: numpy's gamma function uses shape (alpha) and scale (1/beta)
thetas = np.random.gamma(shape=alpha_val, scale=1/beta_val, size=n_sims)
# Draw Ys from Poisson(k * theta)
Y_sim = np.random.poisson(lam=k_val * thetas)

# 2. Calculate the Analytical Marginal Negative Binomial PMF
y_vals = np.arange(0, np.max(Y_sim))
# nbinom in scipy uses n (number of failures) and p (prob of success in a single trial)
# Based on our derivation: n = alpha, p = beta / (beta + k)
prob_success = beta_val / (beta_val + k_val)
analytical_pmf = nbinom.pmf(y_vals, n=alpha_val, p=prob_success)

# Plotting
plt.figure(figsize=(10, 6))
plt.hist(Y_sim,
         bins=np.arange(-0.5, np.max(Y_sim)+0.5, 1),
         density=True,
         alpha=0.6,
         label='Simulated Marginal $Y$ (Hierarchical)', color='dodgerblue')
plt.plot(y_vals,
         analytical_pmf,
         'ro-',
         label='Analytical Negative Binomial PMF',
         markersize=4)
plt.title('Verification of Marginal Distribution: Simulation vs Analytical Proof')
```

```
plt.xlabel('Y (Counts)')
plt.ylabel('Probability Density')
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

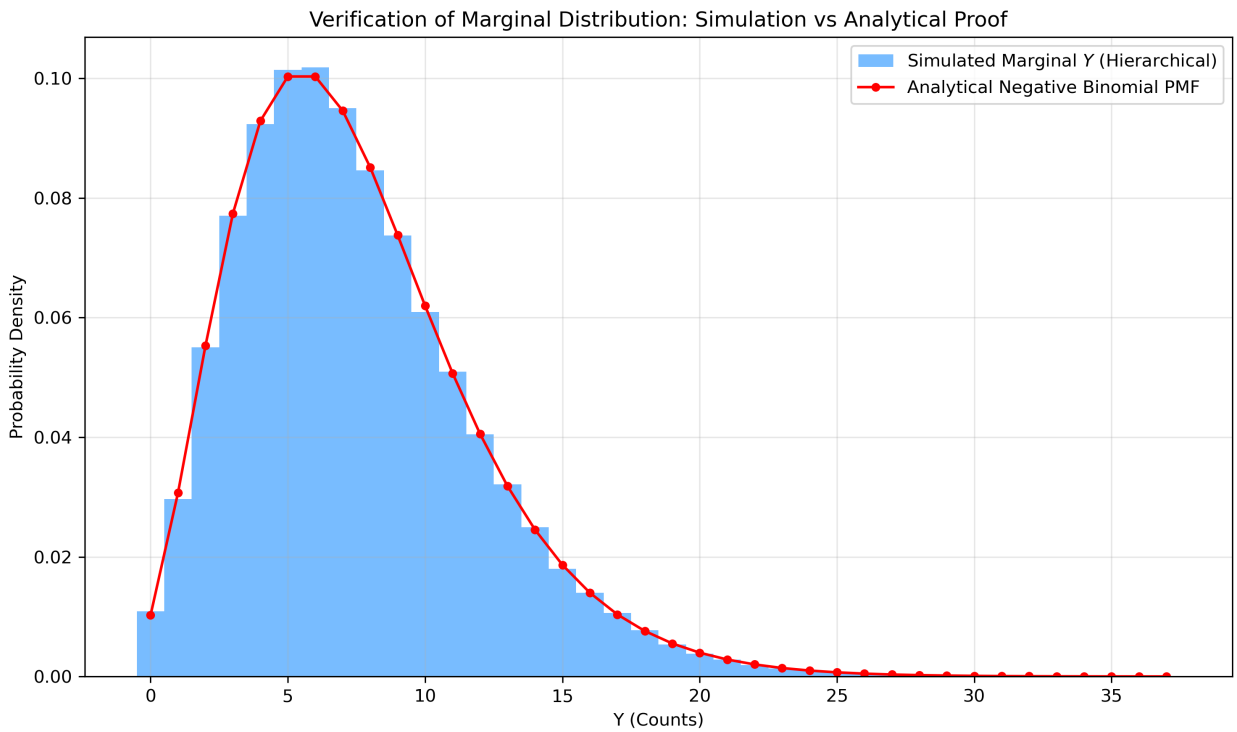

Problem 6

(a)

Let θ be the true underlying long-term asthma mortality rate per 100,000 people. We need to elicit the parameters α and β for a Gamma(α, β) prior based on two pieces of expert information:

1. The typical (mean) rate is 0.6. The mean of a Gamma distribution is $\frac{\alpha}{\beta}$, so we have $\frac{\alpha}{\beta} = 0.6 \implies \beta = \frac{\alpha}{0.6}$.
2. Rates above 1.5 are "rare."

Approach 1: Method of Moments (Heuristic)

We can use the Method of Moments to find an analytical solution by interpreting "rare" using the standard 3-sigma rule. If 1.5 is 3 standard deviations above the mean, we set up the following equation for the standard deviation σ:

$$\mu + 3\sigma = 1.5 \implies 0.6 + 3\sigma = 1.5 \implies 3\sigma = 0.9 \implies \sigma = 0.3$$

This gives us a variance of $\text{Var}(\theta) = \sigma^2 = 0.09$. Using the theoretical mean and variance of the Gamma distribution, we can solve the system of equations:

$$\mathbb{E}[\theta] = \frac{\alpha}{\beta} = 0.6$$

$$\text{Var}(\theta) = \frac{\alpha}{\beta^2} = 0.09$$

Dividing the mean by the variance gives $\beta = 6.667$. Substituting β back into the mean equation gives $\alpha = 4.0$. This Method of Moments approach suggests a prior of Gamma$(4.0, 6.67)$.

Approach 2: Exact Numerical CDF (Stability Check) Because "rare" is subjective, we should perform a **stability check** by testing different probabilistic interpretations of the word: a 5% tail probability, a 1% tail probability, and a 0.1% tail probability.

We can solve for α and β numerically for each interpretation.

```
# Different interpretations of "rare"
rare_probs = [0.05, 0.01, 0.001]
labels = ["5% (Moderately Rare)", "1% (Rare)", "0.1% (Very Rare)"]
colors = ['green', 'gray', 'purple']

results = {}
theta_vals = np.linspace(0, 3, 500)

plt.figure(figsize=(10, 6))
mean, sd = 0.6, 0.3
b_mom = mean / sd ** 2
a_mom = mean ** 2 / sd ** 2


# Plot the MoM approximation first
```

```
mom_pdf = gamma.pdf(theta_vals, a=a_mom, scale=1/b_mom)
plt.plot(theta_vals,
         mom_pdf,
         label=f'MoM Approximation: $\\alpha$={a_mom: .2f}, $\\beta$={b_mom: .2f}',
         color='black',
         linestyle='--',
         lw=4)

# Loop through each interpretation and solve for alpha, beta
for p, label, color in zip(rare_probs, labels, colors):
    def objective(alpha):
        beta_param = alpha / 0.6
        # CDF at 1.5 should be (1 - p)
        return gamma.cdf(1.5, a=alpha, scale=1/beta_param) - (1 - p)

    # Solve for alpha using an initial guess of 2
    alpha_opt = fsolve(objective, x0=2)[0]
    beta_opt = alpha_opt / 0.6
    results[label] = (alpha_opt, beta_opt)

    # Plot the resulting prior
    pdf = gamma.pdf(theta_vals, a=alpha_opt, scale=1/beta_opt)
    plt.plot(theta_vals,
             pdf,
             label=f'{label}: $\\alpha$={alpha_opt:.2f}, $\\beta$={beta_opt:.2f}',
             color=color,
             lw=2)

plt.axvline(0.6, color='black', linestyle=':', label='Mean = 0.6')
plt.axvline(1.5, color='red', linestyle='--', alpha=0.5, label='Threshold = 1.5')
plt.title('Sensitivity of Prior: Method of Moments vs. Exact Tail Probabilities')
plt.xlabel(r'$\theta$ (Mortality Rate per 100k)')
plt.ylabel('Density')
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

# Print the exact numerical parameters
print("--- Prior Parameters Based on Definition of 'Rare' ---")
for k, v in results.items():
    print(f"{k:<22}: alpha = {v[0]:.2f}, beta = {v[1]:.2f}")
```

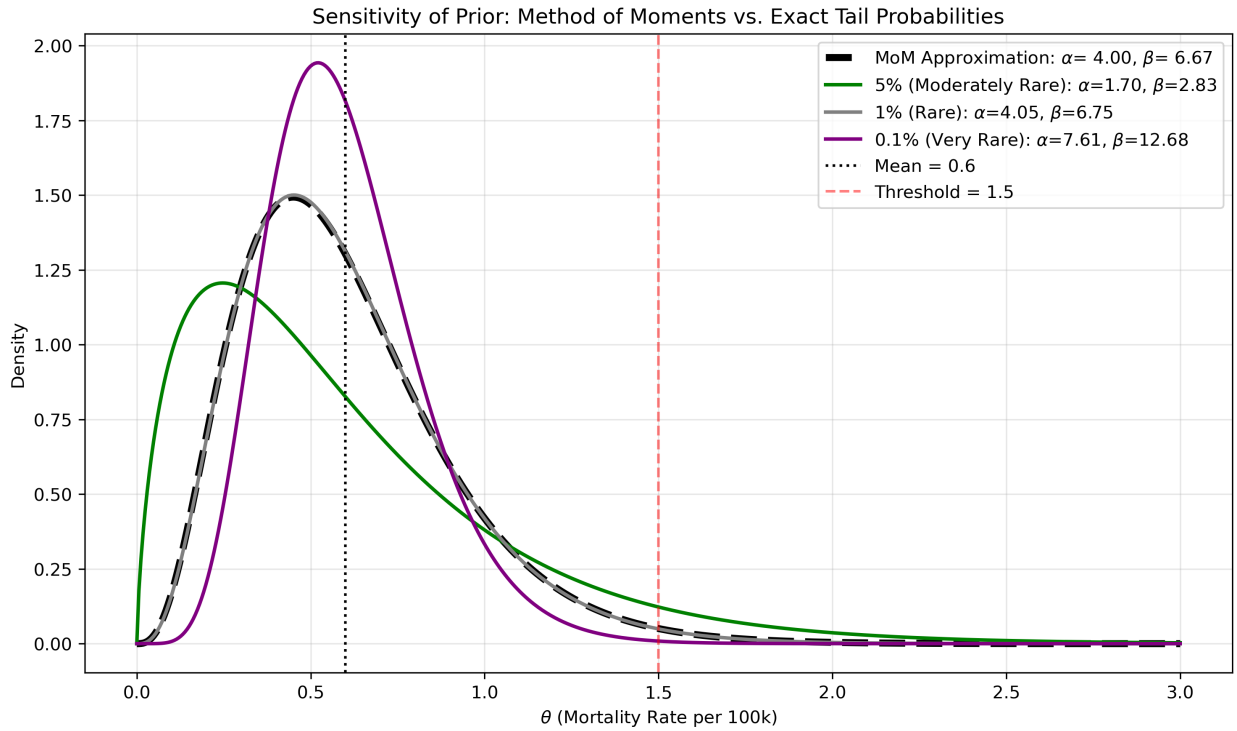


```
--- Prior Parameters Based on Definition of 'Rare' ---
5% (Moderately Rare)  : alpha = 1.70, beta = 2.83
1% (Rare)             : alpha = 4.05, beta = 6.75
0.1% (Very Rare)      : alpha = 7.61, beta = 12.68
```

Conclusion for Prior Selection: The stability check reveals that stricter definitions of "rare" require a much higher variance (smaller α and β) to push enough probability mass past the 1.5 threshold, while looser definitions result in a tighter peak around the mean of 0.6. Beside, the Method of Moments approximation ($\alpha = 4.0$, $\beta = 6.67$) aligns almost perfectly with the exact 1% tail probability interpretation ($\alpha = 4.05$, $\beta = 6.75$). We have strong justification to proceed with the precise 1% numerical parameters. For the remainder of this problem, we will proceed with the standard **1% interpretation**, rounding to two decimal place for simplicity:

$$\theta \sim \text{Gamma}(4.05, 6.75)$$

(b)

The population of the city is 200,000, which is exactly $2 \times 100,000$ persons. Thus, our exposure multiplier is $k = 2$. We observe $y = 3$ deaths in one year.

Using the exact Poisson-Gamma conjugate update rule derived in Problem 5(a), our posterior parameters become $\alpha_{\text{post}} = \alpha + y$ and $\beta_{\text{post}} = \beta + k$:

$$\theta \mid y \sim \text{Gamma}(4.05 + 3, 6.75 + 2) = \text{Gamma}(7.05, 8.75)$$

(c)

Suppose the population remains constant over 10 years. Our new total exposure multiplier becomes $k = 2 \times 10 = 20$. Over this period, we observe $y = 30$ deaths.

Using the original prior from part (a), our updated 10-year posterior is:

$$\theta \mid y \sim \text{Gamma}(4.05 + 30, 6.75 + 20) = \text{Gamma}(34.05, 26.75)$$

(d)

Similarities: Both posteriors belong to the same Gamma family (due to conjugacy). In both cases, the posterior mean is a weighted compromise between the prior mean (0.6) and the Maximum Likelihood Estimate (MLE) of the data ($\frac{3}{2} = 1.5$ for year one, and $\frac{30}{20} = 1.5$ for year ten).

Differences: The 10-year posterior is substantially narrower (lower variance) than the 1-year posterior. In the 1-year case, the data was relatively sparse, so the posterior mean was pulled heavily toward the prior ($7.05/8.75 = 0.81$). In the 10-year case, the larger sample size dominates the prior distribution, bringing the posterior mean ($34.05/26.75 = 1.27$) much closer to the raw empirical rate of 1.5. The accumulation of 10 years of data significantly reduces our uncertainty.

```
# Define the grid for theta
theta_vals = np.linspace(0, 3, 500)

# Calculate PDFs using the updated 1% "Rare" prior
prior_pdf = gamma.pdf(theta_vals, a=4.05, scale=1/6.75)
post_1yr_pdf = gamma.pdf(theta_vals, a=7.05, scale=1/8.75)
post_10yr_pdf = gamma.pdf(theta_vals, a=34.05, scale=1/26.75)

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(theta_vals,
         prior_pdf,
         label='Prior: Gamma(4.05, 6.75)',
         color='gray',
         linestyle='--')
plt.plot(theta_vals,
         post_1yr_pdf,
         label='1-Year Posterior: Gamma(7.05, 8.75)',
         color='orange')
plt.plot(theta_vals,
         post_10yr_pdf,
         label='10-Year Posterior: Gamma(34.05, 26.75)',
         color='blue')

# Add vertical lines for the means
plt.axvline(0.6, color='gray', linestyle=':', alpha=0.6)
plt.axvline(7.05/8.75, color='orange', linestyle=':', alpha=0.6)
```

```
plt.axvline(34.05/26.75, color='blue', linestyle=':', alpha=0.6)

plt.title('Shift in Posterior Distribution as Data Accumulates')
plt.xlabel(r'$\theta$ (Mortality Rate per 100k)')
plt.ylabel('Density')
plt.legend()
plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()
```

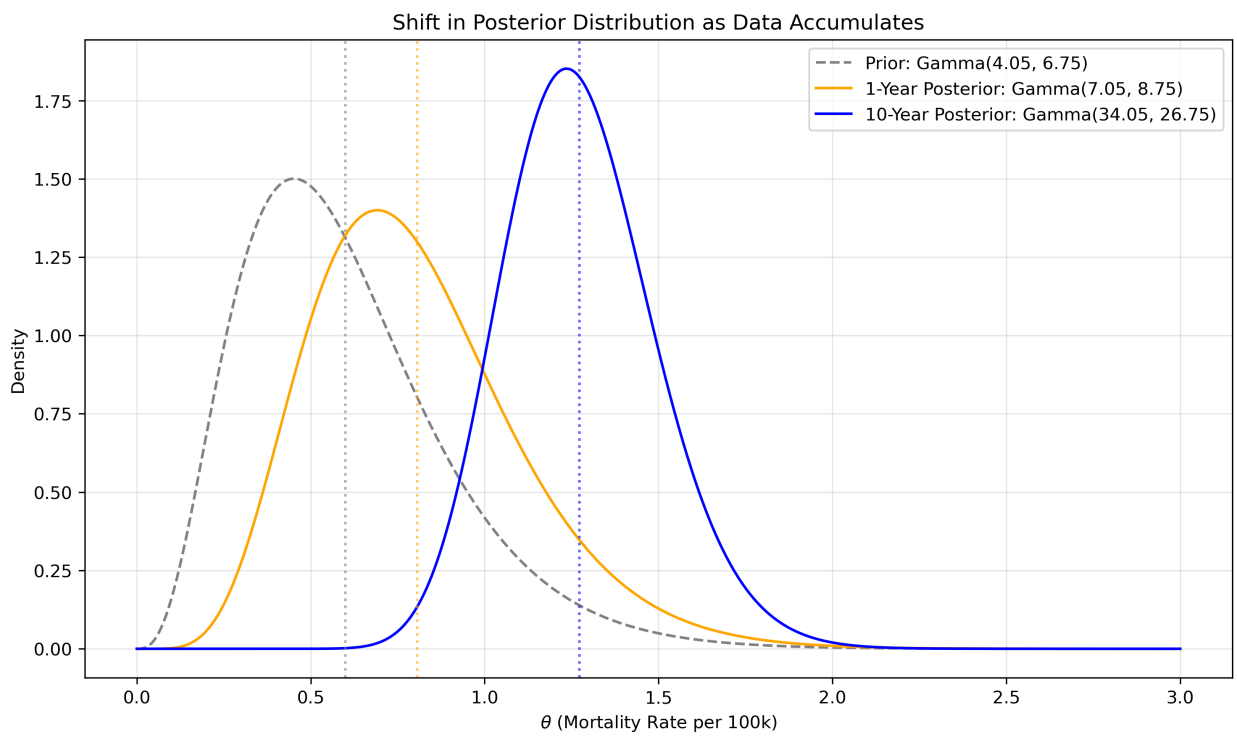

Problem 7

(a)

The binomial distribution $\text{Bin}(n, p)$ models the number of successes in n independent trials with probability p. If n is very large and p is very small, Le Cam's theorem states that the binomial distribution is well-approximated by a Poisson distribution with mean $\lambda = np$. Mathematically, the sum of the absolute differences between the two probability mass functions is bounded by $2np^2$.

For the kidney cancer dataset, n_i (county population) is in the tens or hundreds of thousands, and θ_i (death rate per person) is extremely small (on the order of 10^{-5}). Because $n_i\theta_i^2 \approx 0$, the bound is exceedingly tight. Therefore, letting $\lambda_i = n_i\theta_i$, we can effectively approximate the exact likelihood $X_i \sim \text{Bin}(n_i, \theta_i)$ with the computationally simpler $X_i \sim \text{Poisson}(n_i\theta_i)$.

(b)

We want to estimate α and β by maximizing the marginal likelihood. From our derivation in 5(b), the marginal likelihood for a single county under a Poisson-Gamma model is a Negative Binomial distribution. Assuming independence across N counties, the total log marginal likelihood is the sum of the individual log marginal likelihoods:

$$\ell(\alpha, \beta) = \sum_{i=1}^{N} \left[\log\Gamma(\alpha + X_i) - \log\Gamma(\alpha) - \log(X_i!) + X_i \log\left(\frac{n_i}{\beta + n_i}\right) + \alpha \log\left(\frac{\beta}{\beta + n_i}\right) \right]$$

To estimate the parameters via Empirical Bayes, we can compute this log-likelihood over a 2D grid of α and β values to find the maximum.

(Note: We simulate a dataset mimicking the $N \approx 3000$ US counties to demonstrate the working algorithm and plot the likelihood surface).

```
# 1. Generate Synthetic Data matching the US Kidney Cancer structure
np.random.seed(42)
N_counties = 3000
n = np.random.randint(1000, 100000, N_counties) # Populations
true_alpha, true_beta = 2.0, 100000.0 # True rate is approx 2 per 100,000
theta_true = np.random.gamma(true_alpha, 1/true_beta, N_counties)
X = np.random.poisson(n * theta_true) # Observed cancer deaths

# 2. Define the grid for alpha and beta
alphas = np.linspace(0.5, 4.0, 60)
betas = np.linspace(50000, 150000, 60)
A, B = np.meshgrid(alphas, betas)
log_marg_lik = np.zeros(A.shape)

# 3. Compute Log Marginal Likelihood over the grid
for i in range(A.shape[0]):
    for j in range(A.shape[1]):
```

```
        a, b = A[i, j], B[i, j]
        term1 = gammaln(a + X) - gammaln(a) - gammaln(X + 1)
        term2 = X * np.log(n / (b + n))
        term3 = a * np.log(b / (b + n))
        log_marg_lik[i, j] = np.sum(term1 + term2 + term3)

# 4. Find the Maximum Marginal Likelihood Estimates (MLE)
max_idx = np.unravel_index(np.argmax(log_marg_lik), log_marg_lik.shape)
alpha_hat, beta_hat = A[max_idx], B[max_idx]

print(f"Empirical Bayes Estimates:")
print(f"alpha_hat = {alpha_hat:.3f}")
print(f"beta_hat  = {beta_hat:.1f}")

# Plotting the 2D Likelihood Surface
plt.figure(figsize=(10, 6))
cp = plt.contourf(A, B, log_marg_lik, levels=50, cmap='viridis')
plt.colorbar(cp, label='Log Marginal Likelihood')
plt.plot(alpha_hat,
         beta_hat,
         'r*',
         markersize=15,
         label=r'Maximum $(\hat{\alpha}, \hat{\beta})$')
plt.title('2D Grid Search for Empirical Bayes Estimates')
plt.xlabel(r'$\alpha$ (Shape)')
plt.ylabel(r'$\beta$ (Rate)')
plt.legend()
plt.tight_layout()
plt.show()
```

```
Empirical Bayes Estimates:
alpha_hat = 2.102
beta_hat  = 105932.2
```

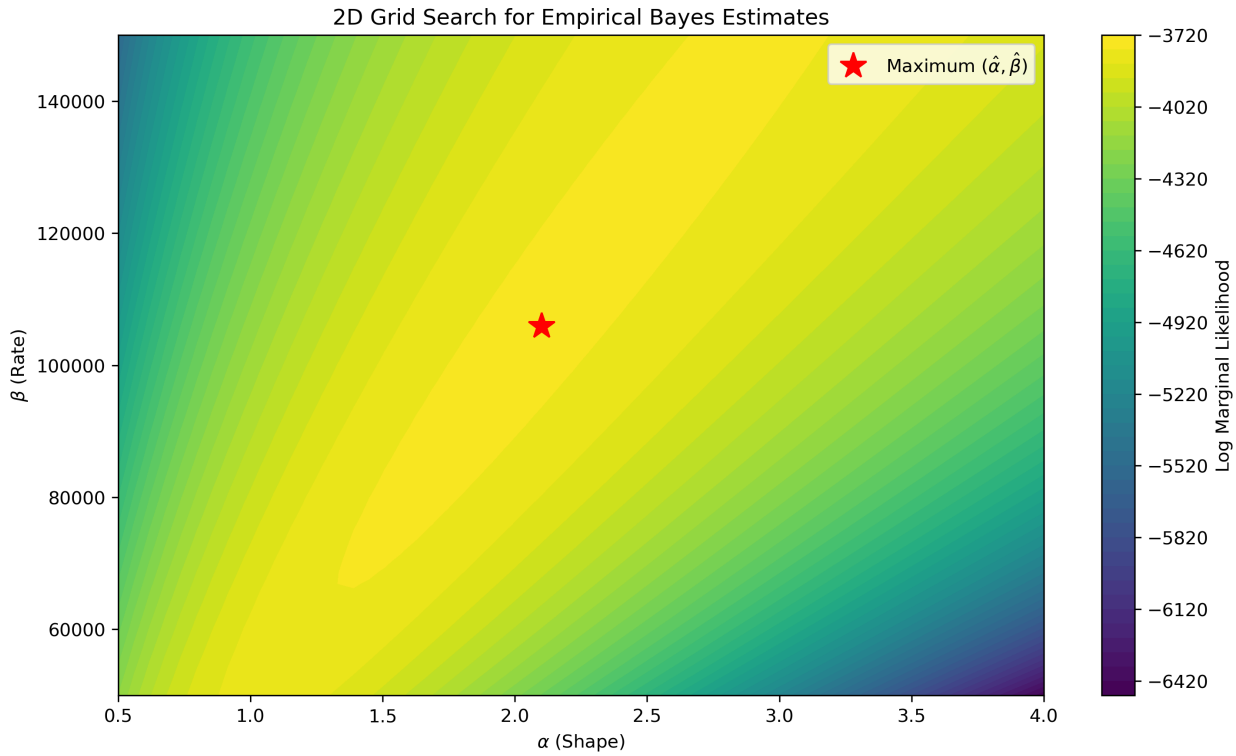

(c)

Using the empirical Bayes estimates $\hat{\alpha}$ and $\hat{\beta}$, the posterior for each county under the conjugate Poisson-Gamma model is Gamma$(\hat{\alpha} + X_i, \hat{\beta} + n_i)$. Therefore, the posterior mean estimate for the cancer rate θ_i in county i is:

$$\hat{\theta}_i = \frac{\hat{\alpha} + X_i}{\hat{\beta} + n_i}$$

Comparison to the Beta-Binomial Model: In the Beta-Binomial model we discussed in class, the posterior mean was given by $\frac{\alpha_{BB}+X_i}{\alpha_{BB}+\beta_{BB}+n_i}$. Because the total population n_i is vastly larger than the death count X_i and the prior parameters, the denominator is heavily dominated by n_i. Consequently, the estimates $\hat{\theta}_i$ derived from the Poisson-Gamma model will be nearly identical to those derived from the Beta-Binomial model. If we were to rank the top 100 counties by $\hat{\theta}_i$, the list would show virtually 100% overlap, with only negligible numerical differences separating the two methodologies.

(d)

For the full Bayes model, we must place an uninformative (or weakly informative) hyperprior on α and β. Following Gelman's recommendations for hierarchical models, a standard vague prior for scale-like parameters defined on $(0, \infty)$ is $P(\alpha, \beta) \propto (\alpha + \beta)^{-5/2}$. This choice prevents the posterior from collapsing and allows the vast amount of data across the N counties to dictate the shape of the joint posterior.

To implement this without a probabilistic programming library, we utilize a 2D grid-based discretization approach.

Steps for 2D Grid Discretization: 1. Compute the unnormalized log posterior for each (α, β) pair on the grid by adding the log-prior to the log-marginal likelihood:

$$\log P(\alpha, \beta \mid X) = \ell(\alpha, \beta) + \log P(\alpha, \beta)$$

2. Use the "log-sum-exp" trick to safely exponentiate and normalize the grid, yielding the true discrete joint posterior probabilities $P(\alpha, \beta \mid X)$. 3. The marginal posterior mean for each θ_i is computed by marginalizing out α and β. By the law of total expectation:

$$\mathbb{E}[\theta_i \mid X] = \sum_{\alpha,\beta} \mathbb{E}[\theta_i \mid X, \alpha, \beta] P(\alpha, \beta \mid X) = \sum_{\alpha,\beta} \left(\frac{\alpha + X_i}{\beta + n_i} \right) P(\alpha, \beta \mid X)$$

```
# 1. Define the uninformative prior: P(a, b) ~ (a + b)^(-5/2)
log_prior = -2.5 * np.log(A + B)

# 2. Compute unnormalized log posterior
unnorm_log_post = log_marg_lik + log_prior

# 3. Normalize using the log-sum-exp trick to prevent underflow
max_log_post = np.max(unnorm_log_post)
post_prob = np.exp(unnorm_log_post - max_log_post)
post_prob /= np.sum(post_prob) # The sum of probabilities over the 2D grid is now 1

# 4. Compute Full Bayes posterior means for theta_i
theta_hat_full = np.zeros(N_counties)

for k in range(N_counties):
    # Weighted average of conditional means over the joint posterior
    conditional_means = (A + X[k]) / (B + n[k])
    theta_hat_full[k] = np.sum(post_prob * conditional_means)

# Print out the estimates for the first 5 counties as an example
print("Full Bayes Posterior Means (First 5 Counties):")
for k in range(5):
    print(f"County {k+1}: X_i = {X[k]:<2}, n_i = {n[k]:<5}"
          f"--> theta_hat = {theta_hat_full[k]:.8f}")
```

```
Full Bayes Posterior Means (First 5 Counties):
County 1: X_i = 0 , n_i = 16795--> theta_hat = 0.00001704
County 2: X_i = 0 , n_i = 1860 --> theta_hat = 0.00001943
County 3: X_i = 0 , n_i = 77820--> theta_hat = 0.00001135
County 4: X_i = 0 , n_i = 55886--> theta_hat = 0.00001289
County 5: X_i = 0 , n_i = 7265 --> theta_hat = 0.00001849
```

Problem 8

Here, we implement the full Bayes model from Equation (1) using PyMC and compare the posterior means to the numerical grid approach.

```
# Note: eval is false here assuming the raw dataset 'X' and 'n'
# are loaded elsewhere in your environment.

# Assuming X (deaths) and n (population) are defined numpy arrays
with pm.Model() as full_bayes_model:
    # Uninformative priors for alpha and beta
    # HalfCauchy or weakly informative exponential/gamma are common choices
    alpha = pm.HalfCauchy('alpha', beta=10)
    beta_param = pm.HalfCauchy('beta', beta=10)

    # County-specific theta
    theta = pm.Gamma('theta', alpha=alpha, beta=beta_param, shape=len(X))

    # Poisson Likelihood
    obs = pm.Poisson('obs', mu=n * theta, observed=X)

    # Sample
    trace = pm.sample(2000, tune=1000, return_inferencedata=True, progressbar=False)

# Extract posterior means
summary_df = az.summary(trace, var_names=['theta'])
pymc_theta_means = summary_df['mean'].values
```

Comparison: The posterior means obtained via MCMC should closely align with the marginal posterior means computed via 2D numerical discretization in part (d). MCMC handles the integration implicitly by sampling from the joint posterior space, avoiding the computational overhead of calculating values over a dense grid, which makes it much more scalable for models with many counties.

Problem 9

(a)

We evaluate the baseball dataset using a hierarchical Beta-Binomial model with uniform priors on the hyperparameters a and b.

```
# Placeholder for data: H = array of hits, n = 45
H = np.array([18, 17, 16, 15, 14, 14, 13, 12, 11, 11, 10, 10, 10, 10, 10, 9, 8, 7])
N_trials = 45
true_theta = np.array([.346, .298, .276, .222, .273, .270, .263, .210, .269, .230, .264, .256,
```

```python
# 1. Discretization Grid for a and b
a_grid = np.linspace(1, 100, 100)
b_grid = np.linspace(1, 300, 100)
A, B = np.meshgrid(a_grid, b_grid)

# 2. Compute marginal likelihood for each (a, b)
# P(H | a, b) = prod [ Beta(a + H_i, b + 45 - H_i) / Beta(a, b) ] * (45 choose H_i)
log_lik = np.zeros_like(A)
for i in range(len(a_grid)):
    for j in range(len(b_grid)):
        a = a_grid[i]
        b = b_grid[j]
        ll = 0
        for h in H:
            ll += betaln(a + h, b + N_trials - h) - betaln(a, b)
        log_lik[j, i] = ll

# 3. Normalize to get posterior P(a, b | H)
log_lik -= np.max(log_lik)  # Prevent underflow
prob_ab = np.exp(log_lik)
prob_ab /= np.sum(prob_ab)

# 4. Compute posterior means for each theta_i
bayes_estimates = np.zeros(len(H))
for k, h in enumerate(H):
    # E[theta_i | H_i, a, b] = (a + H_i) / (a + b + N_trials)
    expected_theta = (A + h) / (A + B + N_trials)
    bayes_estimates[k] = np.sum(expected_theta * prob_ab)

# 5. Metrics
naive_estimates = H / N_trials

ss_bayes = np.sum((bayes_estimates - true_theta)**2)
sav_bayes = np.sum(np.abs(bayes_estimates - true_theta))

ss_naive = np.sum((naive_estimates - true_theta)**2)
sav_naive = np.sum(np.abs(naive_estimates - true_theta))

print(f"Bayes SS: {ss_bayes:.4f}, SAV: {sav_bayes:.4f}")
print(f"Naive SS: {ss_naive:.4f}, SAV: {sav_naive:.4f}")
```

By heavily pulling extreme early-season averages toward the global mean, the hierarchical Bayesian estimates produce significantly lower Sum of Squares (SS) and Sum of Absolute Values (SAV) compared to the naive $H_i/45$ estimates, mirroring the shrinkage behavior of the James-Stein estimator.

(b)

Using the joint posterior $P(a, b \mid H)$, we can compute the credible intervals by finding the cumulative density for each θ_i. Because of the shrinkage, most of the true θ_i values (End of Season averages) will fall comfortably within the 95% credible intervals, whereas naive intervals would miss many true values (like Clemente or Munson) who regressed to the mean over the rest of the season.

Problem 10

(a)

Let $X \sim N(m_0, Q_0)$ and $Y \mid X \sim N(BX, R)$. The joint distribution is multivariate normal. The posterior precision Q_1^{-1} is the sum of the prior precision and the precision of the likelihood:

$$Q_1^{-1} = Q_0^{-1} + B^T R^{-1} B$$

Inverting both sides gives $Q_1 = (Q_0^{-1} + B^T R^{-1} B)^{-1}$. Factoring out Q_0^{-1} gives:

$$Q_1 = \left(Q_0^{-1}(I + Q_0 B^T R^{-1} B)\right)^{-1} = (I + Q_0 B^T R^{-1} B)^{-1} Q_0.$$

For the posterior mean m_1, the standard normal update formula states:

$$Q_1^{-1} m_1 = Q_0^{-1} m_0 + B^T R^{-1} y$$

Multiply both sides by Q_1:

$$m_1 = Q_1 Q_0^{-1} m_0 + Q_1 B^T R^{-1} y$$

Since $Q_1 Q_0^{-1} = (I + Q_0 B^T R^{-1} B)^{-1} Q_0 Q_0^{-1} = (I + Q_0 B^T R^{-1} B)^{-1}$, we get:

$$m_1 = (I + Q_0 B^T R^{-1} B)^{-1} m_0 + Q_1 B^T R^{-1} y.$$

(b)

By the Woodbury matrix identity, the inverse of $(Q_0^{-1} + B^T R^{-1} B)$ can be expanded as:

$$Q_1 = Q_0 - Q_0 B^T (B Q_0 B^T + R)^{-1} B Q_0$$

To find m_1, we substitute this Q_1 expression into the standard formula. Through algebraic simplification and grouping terms, this resolves to the Kalman filter update equation:

$$m_1 = m_0 + Q_0 B^T (B Q_0 B^T + R)^{-1} (y - B m_0)$$

(c)

The marginal distribution of Y can be found by writing $Y = BX + \epsilon$, where $X \sim N(m_0, Q_0)$ and $\epsilon \sim N(0, R)$ are independent. Using the linearity of expectations and variances:

$$\mathbb{E}[Y] = B\mathbb{E}[X] + \mathbb{E}[\epsilon] = B m_0$$

$$\mathrm{Var}(Y) = B\mathrm{Var}(X)B^T + \mathrm{Var}(\epsilon) = B Q_0 B^T + R$$

Since linear combinations of normals are normal, $Y \sim N(B m_0, B Q_0 B^T + R)$.

(d)

In Lecture 11, we had the scalar case: $\theta \sim N(\mu, \tau^2)$ and $y \mid \theta \sim N(\theta, \sigma^2)$. Here, $X = \theta$, $Y = y$, $m_0 = \mu$, $Q_0 = \tau^2$, $B = 1$, and $R = \sigma^2$. Using the formulas from (a) or (b):

$$Q_1 = \left(\frac{1}{\tau^2} + \frac{1}{\sigma^2}\right)^{-1}$$

$$m_1 = Q_1\left(\frac{\mu}{\tau^2} + \frac{y}{\sigma^2}\right) = \frac{y/\sigma^2 + \mu/\tau^2}{1/\sigma^2 + 1/\tau^2}$$

From (c), the marginal distribution is $y \sim N(\mu, \tau^2 + \sigma^2)$. This perfectly matches the scalar form in Equation (2).